

A Comprehensive Guide to Pandas

Your Name

Saturday 28th December, 2024

1 Introduction

Pandas is a powerful and flexible open-source data analysis and manipulation library for Python. Built on top of NumPy, it provides easy-to-use data structures and data analysis tools. The primary data structures in Pandas are Series and DataFrame.

2 Installation

To start using Pandas, you need to install it. You can do this using pip, the package installer for Python. Open your terminal or command prompt and run the following command:

```
pip install pandas
```

3 Basic Data Structures

Pandas offers two main data structures: **Series** and **DataFrame**.

3.1 Series

A Series is a one-dimensional labeled array that can hold any data type. It is similar to a list or an array but with additional features.

3.1.1 Creating a Series

You can create a Series from a list, dictionary, or NumPy array.

```
import pandas as pd

# Creating a Series from a list
data = [10, 20, 30, 40]
series = pd.Series(data)
print("Series from list:")
```

```

print(series)

# Creating a Series from a dictionary
data_dict = {'a': 1, 'b': 2, 'c': 3}
series_dict = pd.Series(data_dict)
print("\nSeries from dictionary:")
print(series_dict)

# Creating a Series from a NumPy array
import numpy as np
numpy_array = np.array([1, 2, 3, 4])
series_from_array = pd.Series(numpy_array)
print("\nSeries from NumPy array:")
print(series_from_array)

```

3.1.2 Output

The output of the above code will be:

Series from list:

```

0    10
1    20
2    30
3    40
dtype: int64

```

Series from dictionary:

```

a     1
b     2
c     3
dtype: int64

```

Series from NumPy array:

```

0     1
1     2
2     3
3     4
dtype: int64

```

3.2 DataFrame

A DataFrame is a two-dimensional labeled data structure with columns of potentially different types. It is similar to a spreadsheet or SQL table.

3.2.1 Creating a DataFrame

You can create a DataFrame from a dictionary, a list of lists, or a NumPy array.

```

# Creating a DataFrame from a dictionary
data = {
    'Name': ['Alice', 'Bob', 'Charlie'],
    'Age': [25, 30, 35],
    'City': ['New_York', 'Los_Angeles', 'Chicago']
}
df = pd.DataFrame(data)
print("\nDataFrame:")
print(df)

# Creating a DataFrame from a list of lists
data_list = [['Alice', 25, 'New_York'], ['Bob', 30, 'Los_Angeles'], ['Charlie', 35, 'Chicago']]
df_from_list = pd.DataFrame(data_list, columns=['Name', 'Age', 'City'])
print("\nDataFrame_from_list_of_lists:")
print(df_from_list)

```

3.2.2 Output

The output of the above code will be:

```

DataFrame:
   Name  Age  City
0  Alice  25  New York
1   Bob   30 Los Angeles
2 Charlie  35   Chicago

```

```

DataFrame from list of lists:
   Name  Age  City
0  Alice  25  New York
1   Bob   30 Los Angeles
2 Charlie  35   Chicago

```

3.3 Example DataFrame

The following table represents a sample DataFrame created from the previous code:

| Table 1: Sample DataFrame | | |
|---------------------------|-----|-------------|
| Name | Age | City |
| Alice | 25 | New York |
| Bob | 30 | Los Angeles |
| Charlie | 35 | Chicago |

4 Basic Operations

Pandas provides various operations to manipulate and analyze data.

4.1 Accessing Data

You can access rows and columns using labels or indices.

```
# Accessing a column
print("\nAccessing a column:")
print(df['Name'])

# Accessing multiple columns
print("\nAccessing multiple columns:")
print(df[['Name', 'Age']])

# Accessing a row by index
print("\nAccessing a row by index:")
print(df.loc[1]) # Accessing the second row

# Accessing a row by index position
print("\nAccessing a row by index position:")
print(df.iloc[1]) # Accessing the second row
```

4.1.1 Output

The output of the above code will be:

Accessing a column:

```
0    Alice
1     Bob
2   Charlie
Name: Name, dtype: object
```

Accessing multiple columns:

```
   Name  Age
0  Alice  25
1   Bob   30
2 Charlie  35
```

Accessing a row by index:

```
Name    Bob
Age      30
City  Los Angeles
Name: 1, dtype: object
```

Accessing a row by index position:

```
Name      Bob
Age       30
City      Los Angeles
Name: 1, dtype: object
```

4.2 Filtering Data

You can filter data based on conditions using boolean indexing.

```
# Filtering rows where Age is greater than 28
filtered_df = df[df['Age'] > 28]
print("\nFiltered DataFrame (Age > 28):")
print(filtered_df)

# Filtering with multiple conditions
filtered_multiple = df[(df['Age'] > 28) & (df['City']
    == 'Los Angeles')]
print("\nFiltered DataFrame (Age > 28 and City is Los Angeles):")
print(filtered_multiple)
```

4.2.1 Output

The output of the above code will be:

Filtered DataFrame (Age > 28):

```
   Name  Age  City
1   Bob   30 Los Angeles
2  Charlie 35   Chicago
```

Filtered DataFrame (Age > 28 and City is Los Angeles):

```
   Name  Age  City
1  Bob   30 Los Angeles
```

4.3 Adding and Modifying Columns

You can easily add new columns or modify existing ones.

```
# Adding a new column
df['Salary'] = [70000, 80000, 90000]
print("\nDataFrame after adding Salary column:")
print(df)

# Modifying an existing column
df['Age'] = df['Age'] + 1 # Incrementing age by 1
print("\nDataFrame after modifying Age column:")
print(df)
```

```
# Using apply to modify a column
df['Salary'] = df['Salary'].apply(lambda x: x * 1.1)
# Giving a 10% raise
print("\nDataFrame after applying a function to Salary
      column:")
print(df)
```

4.3.1 Output

The output of the above code will be:

DataFrame after adding Salary column:

| | Name | Age | City | Salary |
|---|---------|-----|-------------|--------|
| 0 | Alice | 25 | New York | 70000 |
| 1 | Bob | 30 | Los Angeles | 80000 |
| 2 | Charlie | 35 | Chicago | 90000 |

DataFrame after modifying Age column:

| | Name | Age | City | Salary |
|---|---------|-----|-------------|--------|
| 0 | Alice | 26 | New York | 70000 |
| 1 | Bob | 31 | Los Angeles | 80000 |
| 2 | Charlie | 36 | Chicago | 90000 |

DataFrame after applying a function to Salary column:

| | Name | Age | City | Salary |
|---|---------|-----|-------------|---------|
| 0 | Alice | 26 | New York | 77000.0 |
| 1 | Bob | 31 | Los Angeles | 88000.0 |
| 2 | Charlie | 36 | Chicago | 99000.0 |

4.4 Example DataFrame After Modifications

The following table represents the DataFrame after adding the Salary column and modifying the Age column:

| Table 2: Modified DataFrame | | | |
|-----------------------------|-----|-------------|---------|
| Name | Age | City | Salary |
| Alice | 26 | New York | 77000.0 |
| Bob | 31 | Los Angeles | 88000.0 |
| Charlie | 36 | Chicago | 99000.0 |

5 Data Manipulation

Pandas allows for a variety of data manipulation techniques.

5.1 Grouping Data

You can group data and perform aggregate functions such as sum, mean, etc.

```
# Grouping by City and calculating average Salary
grouped_df = df.groupby('City')['Salary'].mean().
    reset_index()
print("\nGrouped DataFrame (Average Salary by City):")
print(grouped_df)

# Grouping by multiple columns
grouped_multiple = df.groupby(['City', 'Age']).size().
    reset_index(name='Counts')
print("\nGrouped DataFrame by City and Age:")
print(grouped_multiple)
```

5.1.1 Output

The output of the above code will be:

Grouped DataFrame (Average Salary by City):

| | City | Salary |
|---|-------------|---------|
| 0 | Chicago | 99000.0 |
| 1 | Los Angeles | 88000.0 |
| 2 | New York | 77000.0 |

5.2 Sorting Data

You can sort a DataFrame by one or multiple columns.

```
# Sorting by Age
sorted_df = df.sort_values(by='Age', ascending=False)
print("\nSorted DataFrame by Age (Descending):")
print(sorted_df)

# Sorting by multiple columns
sorted_multiple = df.sort_values(by=['City', 'Age'],
    ascending=[True, False])
print("\nSorted DataFrame by City (Ascending) and Age (Descending):")
print(sorted_multiple)
```

5.2.1 Output

The output of the above code will be:

Sorted DataFrame by Age (Descending):

| | Name | Age | City | Salary |
|--|------|-----|------|--------|
|--|------|-----|------|--------|

```

2 Charlie 36 Chicago 99000.0
1 Bob 31 Los Angeles 88000.0
0 Alice 26 New York 77000.0

```

Sorted DataFrame by City (Ascending) and Age (Descending):

```

      Name Age      City Salary
2 Charlie 36 Chicago 99000.0
1 Bob 31 Los Angeles 88000.0
0 Alice 26 New York 77000.0

```

6 Handling Missing Data

Dealing with missing data is an essential part of data analysis.

6.1 Identifying Missing Data

You can check for missing values using the `isnull()` method.

```

# Introducing some missing values
df.loc[1, 'City'] = None
print("\nDataFrame with missing values:")
print(df)

# Identifying missing values
print("\nMissing values in DataFrame:")
print(df.isnull().sum()) % Count of missing values in
                        each column

```

6.1.1 Output

The output of the above code will be:

DataFrame with missing values:

```

      Name      Age      City Salary
0  Alice 26.000000 New York 77000.0
1   Bob 31.000000     None 88000.0
2 Charlie 36.000000 Chicago 99000.0

```

Missing values in DataFrame:

```

Name      0
Age       0
City      1
Salary    0
dtype: int64

```


6.2 Handling Missing Data

You can either drop rows with missing values or fill them with a specific value.

```
# Dropping rows with missing values
df_dropped = df.dropna()
print("\nDataFrame after dropping missing values:")
print(df_dropped)

# Filling missing values
df_filled = df.fillna('Unknown')
print("\nDataFrame after filling missing values:")
print(df_filled)
```

6.2.1 Output

The output of the above code will be:

DataFrame after dropping missing values:

| | Name | Age | City | Salary |
|---|---------|-----------|----------|---------|
| 0 | Alice | 26.000000 | New York | 77000.0 |
| 2 | Charlie | 36.000000 | Chicago | 99000.0 |

DataFrame after filling missing values:

| | Name | Age | City | Salary |
|---|---------|-----------|----------|---------|
| 0 | Alice | 26.000000 | New York | 77000.0 |
| 1 | Bob | 31.000000 | Unknown | 88000.0 |
| 2 | Charlie | 36.000000 | Chicago | 99000.0 |

7 Data Type Conversion

Pandas allows you to check and convert data types of DataFrame columns.

7.1 Checking Data Types

You can check the data types of each column using the `dtypes` attribute.

```
print("\nData Types:")
print(df.dtypes) % Checking data types of columns
```

7.1.1 Output

The output of the above code will be:

Data Types:

| | |
|------|--------|
| Name | object |
| Age | object |

```
City      object
Salary    float64
dtype: object
```

7.2 Converting Data Types

You can convert the data types of columns using the `astype()` method.

```
# Converting Age to string
df['Age'] = df['Age'].astype(str)
print("\nDataFrame after converting Age to string:")
print(df)
print("Data Types after conversion:")
print(df.dtypes)

# Converting Salary to float
df['Salary'] = df['Salary'].astype(float)
print("\nDataFrame after converting Salary to float:")
print(df)
print("Data Types after conversion:")
print(df.dtypes)
```

7.2.1 Output

The output of the above code will be:

DataFrame after converting Age to string:

| | Name | Age | City | Salary |
|---|---------|-----|----------|---------|
| 0 | Alice | 26 | New York | 77000.0 |
| 1 | Bob | 31 | Unknown | 88000.0 |
| 2 | Charlie | 36 | Chicago | 99000.0 |

Data Types after conversion:

```
Name      object
Age       object
City      object
Salary    float64
dtype: object
```

DataFrame after converting Salary to float:

| | Name | Age | City | Salary |
|---|---------|-----|----------|---------|
| 0 | Alice | 26 | New York | 77000.0 |
| 1 | Bob | 31 | Unknown | 88000.0 |
| 2 | Charlie | 36 | Chicago | 99000.0 |

Data Types after conversion:

```
Name      object
Age       object
City      object
```

```
Salary    float64
dtype: object
```

8 Best Practices and Tips

- **Use Descriptive Column Names:** Use clear and descriptive names for your columns to make your DataFrame easier to understand.
- **Check Data Types:** Always check the data types of your DataFrame to ensure they are appropriate for your analysis.
- **Handle Missing Data Early:** Identify and handle missing data as early as possible to avoid complications later in your analysis.
- **Use Vectorized Operations:** Prefer vectorized operations (like applying functions to entire columns) over loops for better performance.
- **Document Your Code:** Comment on your code to explain complex operations or decisions you make during data manipulation.
- **Explore Data:** Use methods like `head()`, `tail()`, and `describe()` to explore your DataFrame and understand its structure.
- **Be Cautious with Type Conversions:** Ensure that the data being converted is compatible with the target type to avoid errors.

9 Comprehensive Example: Data Analysis Workflow

In this section, we will apply everything we have learned using a sample dataset. We will perform the following steps:

1. Load the dataset.
2. Clean the data (handling missing values).
3. Perform exploratory data analysis (EDA).
4. Manipulate the data (add and modify columns).
5. Group and summarize the data.
6. Visualize the results.

9.1 Step 1: Load the Dataset

Let's create a sample dataset representing employee records.

```
# Sample dataset
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eva',
            , None],
    'Age': [25, 30, 35, None, 22, 28],
    'City': ['New_York', 'Los_Angeles', 'Chicago', 'Miami', None, 'Seattle'],
}
```

```

        'Salary': [70000, 80000, 90000, 100000, 95000,
                    85000]
    }
    df = pd.DataFrame(data)
    print("\nInitial DataFrame:")
    print(df)

```

9.1.1 Output

The output of the above code will be:

Initial DataFrame:

| | Name | Age | City | Salary |
|---|---------|------|-------------|--------|
| 0 | Alice | 25.0 | New York | 70000 |
| 1 | Bob | 30.0 | Los Angeles | 80000 |
| 2 | Charlie | 35.0 | Chicago | 90000 |
| 3 | David | NaN | Miami | 100000 |
| 4 | Eva | 22.0 | None | 95000 |
| 5 | None | 28.0 | Seattle | 85000 |

9.2 Step 2: Clean the Data

We will handle missing values by filling them with appropriate defaults.

```

# Filling missing values
df['Name'].fillna('Unknown', inplace=True)
df['City'].fillna('Unknown', inplace=True)
df['Age'].fillna(df['Age'].mean(), inplace=True) #
Filling with mean age
print("\nDataFrame after cleaning missing values:")
print(df)

```

9.2.1 Output

The output of the above code will be:

DataFrame after cleaning missing values:

| | Name | Age | City | Salary |
|---|---------|-----------|-------------|--------|
| 0 | Alice | 25.000000 | New York | 70000 |
| 1 | Bob | 30.000000 | Los Angeles | 80000 |
| 2 | Charlie | 35.000000 | Chicago | 90000 |
| 3 | David | 28.333333 | Miami | 100000 |
| 4 | Unknown | 22.000000 | Unknown | 95000 |
| 5 | Unknown | 28.000000 | Seattle | 85000 |

9.3 Step 3: Exploratory Data Analysis (EDA)

We will perform some basic EDA to understand the data better.

```
# Descriptive statistics
print("\nDescriptive statistics:")
print(df.describe(include='all'))

# Checking unique values in 'City'
print("\nUnique cities:")
print(df['City'].unique())
```

9.3.1 Output

The output of the above code will be:

Descriptive statistics:

| | Name | Age | City | Salary |
|--------|---------|-----------|----------|---------------|
| count | 6.0 | 6.000000 | 6 | 6.000000 |
| unique | 6.0 | NaN | 5 | NaN |
| top | Unknown | NaN | New York | NaN |
| freq | 2.0 | NaN | 1 | NaN |
| mean | NaN | 28.333333 | NaN | 82500.000000 |
| std | NaN | 4.163334 | NaN | 10954.451580 |
| min | NaN | 22.000000 | NaN | 70000.000000 |
| 25% | NaN | 25.000000 | NaN | 75000.000000 |
| 50% | NaN | 28.000000 | NaN | 82500.000000 |
| 75% | NaN | 30.000000 | NaN | 90000.000000 |
| max | NaN | 35.000000 | NaN | 100000.000000 |

Unique cities:

```
['New York' 'Los Angeles' 'Chicago' 'Miami' 'Unknown' 'Seattle']
```

9.4 Step 4: Manipulate the Data

We will add a new column for the 'Country' and modify the 'Salary' column.

```
# Adding a new column 'Country'
df['Country'] = 'USA'
print("\nDataFrame after adding Country column:")
print(df)

# Modifying the Salary column (giving a 5% bonus)
df['Salary'] = df['Salary'] * 1.05
print("\nDataFrame after applying a 5% bonus to Salary:")
print(df)
```

9.4.1 Output

The output of the above code will be:

DataFrame after adding Country column:

| | Name | Age | City | Salary | Country |
|---|---------|-----------|-------------|--------|---------|
| 0 | Alice | 26.000000 | New York | 73500 | USA |
| 1 | Bob | 31.000000 | Los Angeles | 84000 | USA |
| 2 | Charlie | 36.000000 | Chicago | 94500 | USA |
| 3 | David | 28.333333 | Miami | 105000 | USA |
| 4 | Unknown | 22.000000 | Unknown | 99750 | USA |
| 5 | Unknown | 28.000000 | Seattle | 89250 | USA |

DataFrame after applying a 5% bonus to Salary:

| | Name | Age | City | Salary | Country |
|---|---------|-----------|-------------|----------|---------|
| 0 | Alice | 26.000000 | New York | 77000.0 | USA |
| 1 | Bob | 31.000000 | Los Angeles | 88000.0 | USA |
| 2 | Charlie | 36.000000 | Chicago | 99000.0 | USA |
| 3 | David | 28.333333 | Miami | 105000.0 | USA |
| 4 | Unknown | 22.000000 | Unknown | 99750.0 | USA |
| 5 | Unknown | 28.000000 | Seattle | 89250.0 | USA |

9.5 Step 5: Group and Summarize the Data

We will group the data by 'City' and calculate the average salary.

```
# Grouping by City and calculating average Salary
grouped_salary = df.groupby('City')['Salary'].mean().
    reset_index()
print("\nGrouped DataFrame (Average Salary by City):")
print(grouped_salary)
```

9.5.1 Output

The output of the above code will be:

Grouped DataFrame (Average Salary by City):

| | City | Salary |
|---|-------------|----------|
| 0 | Chicago | 99000.0 |
| 1 | Los Angeles | 84000.0 |
| 2 | Miami | 105000.0 |
| 3 | New York | 77000.0 |
| 4 | Seattle | 89250.0 |
| 5 | Unknown | 99750.0 |

9.6 Step 6: Visualize the Results


We can visualize the average salary by city using a bar chart.

```
import matplotlib.pyplot as plt

# Plotting the average salary by city
plt.figure(figsize=(10, 6))
plt.bar(grouped_salary['City'], grouped_salary['Salary'], color='skyblue')
plt.title('Average Salary by City')
plt.xlabel('City')
plt.ylabel('Average Salary ($)')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

9.7 Workflow Summary

The following flowchart summarizes the data analysis workflow:



10 Conclusion

This comprehensive example demonstrated how to apply the concepts learned throughout this guide using a sample dataset. We covered loading the dataset, cleaning the data, performing exploratory data analysis, manipulating the data, grouping and summarizing the data, and visualizing the results.

By practicing these steps, you will strengthen your understanding of Pandas and gain confidence in using it for data analysis tasks.